

MP3: CPU Scheduling

Cover page

- 徐竣霆 : Trace code, Implementation, Report.
- 江承紘 : Trace code, Implementation, Report.

Trace Code

1-1.New->Ready

Kernel::ExecAll()

Iterate through all the execfile, Exec() them. And set the current thread to Finished(Sleep())(Parent finished after all the children finished).

Kernel::Exec(char *)

Create a new thread and allocate Address Space for the thread. Call Fork().

Thread::Fork(VoidFunctionPtr, void*)

Fork() allows caller(parent) and callee(child) to execute concurrently. Call StackAllocate() to allocate stack for child. Then call ReadyToRun()(Should disable interrupt before, and enable interrupt afterward).

Thread::StackAllocate(VoidFunctionPtr, void*)

Use AllocBoundedArray()(in lib/sysdep.cc) to create a continuous space as stack.
This is a x86 system

```
#ifdef x86
    // the x86 passes the return address on the stack. In order for SWITCH()
    // to go to ThreadRoot when we switch to this thread, the return address
    // used in SWITCH() must be the starting address of ThreadRoot.
    stackTop = stack + StackSize - 4;    // -4 to be on the safe side!
    *((--stackTop) = (int) ThreadRoot;
    *stack = STACK_FENCEPOST;
#endif
```

Since the stack address is from high to low, set the stackTop to the highest(the start space, also where ThreadRoot(implemented in switch.S) locates). Then set stack as 0xdedbeef, to detect stack overflows.

```
#else
    machineState[PCState] = (void*)ThreadRoot;
    machineState[StartupPCState] = (void*)ThreadBegin;
    machineState[InitialPCState] = (void*)func;
    machineState[InitialArgState] = (void*)arg;
    machineState[WhenDonePCState] = (void*)ThreadFinish;
#endif
```

Then do this, setting all the machineState.(for Context Switch).(defined in switch.h)

Scheduler::ReadyToRun(Thread*)

Set the status of this Thread to Ready.

And Append this thread to readyList(FIFO, so it's actually a queue).

1-2.Running->Ready

Machine::Run()

Notice that this routine is re-entrant, so that it can run concurrently. (Detail was explained in MP1).

Interrupt::OneTick()

Two things can call OneTick():

1. Interrupts are re-enabled
2. A user instruction is executed

First advance the simulation time(based on current status(system/user)). Then call CheckIfDue() to check for pending interrupts. (turn off/ turn on interrupts before/after)

Check yieldOnReturn(for context switch). If true, call Yield() to yield the CPU from current Thread to others.

Thread::Yield()

Relinquish the CPU, call FindNextToRun() to find the next Thread to run.

Call ReadyToRun() to put the current Thread back to the readyList.

Call Run() to run the next Thread.

Schedulre::FindNextToRun()

If readyList is empty, return NULL directly.

Otherwise return the front of readyList, and pop_front().

Scheduler::ReadyToRun(Thread*)

As 1-1, set current Thread as ready, and push_back() to put it back to the readyList.

Scheduler::Run(Thread*, bool)

Dispatch the CPU to nextThread, check if currentThread(oldThread) is finished. If so, mark currentThread toBeDestroyed. Otherwise, save the state of the oldThread. Switch the currentThread to nextThread and set the state as RUNNING.

Call SWITCH()(the context switch routine, details explained in 1-6 Switch(Thread*, Thread*)).

Call CheckToBeDestroyed() to delete the marked FINISHED thread.

Restore the state of oldThread(if needed).

1-3.Running->Waiting

SynchConsoleOutput::PutChar(char)

Make sure there is only one thread using I/O device at a time.(lock->Acquire())

Write a character to the console display.

Do synchronization(call P()).

Release the lock(lock->Release()).

Semaphore::P()(in synch.cc (<http://synch.cc>))

To handle synchronization.

While value == 0, means semaphore is not availavble, after append the currentThread in the waiting queue, put the currentThread to sleep. Otherwise, value--(semaphore is available).

List::Append(T)(in list.cc (<http://list.cc>))

Check whether the item is in the list already, if not, append it to the end of the list.

Thread::Sleep(bool)

Relinquish the CPU, since currentThread is blocked waiting on semaphore()(in this path).

Set the status of currentThread as BLOCKED.

If there are no threads in ready queue to run, call kernel->interrupt->Idle() to idle the CPU until next I/O interrupt to turn threads from waiting to ready.

Then call Run() to run the nextThread, it will do context switch. When it's time to switch back, it will restore the state and awake the sleeping thread.

Scheduler::FindNextToRun()

explained in 1-2.

Scheduler::Run(Thread*, bool)

explained in 1-2.

1-4.Waiting->Ready(Note: only need to consider console output as an example)

Semaphore::V()(in synch.cc (<http://synch.cc>))

First, call OneTick()(in Run()). Then, call OneTick(). In OneTick(), call checkIfDue() to check for pending interrupts.

If so, pull the front off the interrupt list and call the interrupt handler(next->callOnInterrupt->CallBack()). The CallBack() defined in callback.h, while the CallBackObj in this path is ConsoleOutput. So call ConsoleOutput::CallBack(). This will further call SynchConsoleInput::CallBack()(callWhenDone). And this will do waitFor()->V().

```

void
Semaphore::V()
{
    DEBUG(dbgTraCode, "In Semaphore::V(), " << kernel->stats->totalTicks);
    Interrupt *interrupt = kernel->interrupt;

    // disable interrupts
    IntStatus oldLevel = interrupt->SetLevel(IntOff);

    if (!queue->IsEmpty()) { // make thread ready.
        kernel->scheduler->ReadyToRun(queue->RemoveFront());
    }
    value++;

    // re-enable interrupts
    (void) interrupt->SetLevel(oldLevel);
}

```

The Semaphore::V() cooperate with Semaphore P() to deal with the synchronization problem.

It will first set the front of Semaphore as Ready, put it back to readyList(ReadyToRun()).

Unlike P(), V() will do value++ here(kind of like wait() and signal() in Chapter 6).

Scheduler::ReadyToRun(Thread*)

As explained above.

1-5.Running->Terminated(Note: start from the Exit system call is called)

ExceptionHandler(ExceptionType) case SC_Exit

```

case SC_Exit:
    DEBUG(dbgAddr, "Program exit\n");
    val=kernel->machine->ReadRegister(4);
    cout << "return value:" << val << endl;
    kernel->currentThread->Finish();
    break;

```

will do kernel->currentThread->Finish() to finish currentThread().

Thread::Finish()

Call Sleep()(After IntOff). Notice that we cannot deallocate the current thread immediately while we are still running on this thread. Instead, we have to tell the scheduler to call the destructor, once we are not running this thread(in other

thread).

Thread::Sleep(bool)

Explained in 1-3. Put the currentThread to sleep and find nextThread to run.

Scheduler::FindNextToRun()

Explained in 1-2.

Scheduler::Run(Thread*, bool)

Explained in 1-2.

The marked FINISHED thread will be delete in this case.

1-6.Ready->Running

Scheduler::FindNextToRun()

Explained in 1-2.

Scheduler::Run(Thread*, bool)

Explained in 1-2.

Set the state of nextThread from ready to running, then do SWITCH().

SWITCH(Thread*, Thread*)

(implemented in switch.S)

```
#ifdef x86

    .text
    .align 2

    .globl ThreadRoot
    .globl _ThreadRoot

/* void ThreadRoot( void )
**
** expects the following registers to be initialized:
**     eax    points to startup function (interrupt enable)
**     edx    contains initial argument to thread function
**     esi    points to thread function
**     edi    point to Thread::Finish()
**
**
```

_ThreadRoot:

ThreadRoot:

```
    pushl    %ebp
    movl     %esp,%ebp
    pushl    InitialArg
    call     *StartupPC
    call     *InitialPC
    call     *WhenDonePC
```

NOT REACHED

```
    movl     %ebp,%esp
    popl     %ebp
    ret
```

/* void SWITCH(thread *t1, thread *t2)

**

** on entry, stack looks like this:

```
**      8(esp)  ->          thread *t2
**      4(esp)  ->          thread *t1
**      (esp)  ->          return address
```

**

** we push the current eax on the stack so that we can use it as
** a pointer to t1, this decrements esp by 4, so when we use it
** to reference stuff on the stack, we add 4 to the offset.

*/

```
.comm    _eax_save,4
```

```
.globl   SWITCH
```

```
.globl   _SWITCH
```

_SWITCH:

SWITCH:

```
    movl     %eax,_eax_save          # save the value of eax
    movl     4(%esp),%eax            # move pointer to t1 into eax
    movl     %ebx,_EBX(%eax)         # save registers
    movl     %ecx,_ECX(%eax)
    movl     %edx,_EDX(%eax)
    movl     %esi,_ESI(%eax)
    movl     %edi,_EDI(%eax)
    movl     %ebp,_EBP(%eax)
    movl     %esp,_ESP(%eax)         # save stack pointer
    movl     _eax_save,%ebx          # get the saved value of eax
    movl     %ebx,_EAX(%eax)         # store it
    movl     0(%esp),%ebx            # get return address from stack
    movl     %ebx,_PC(%eax)          # save it into the pc storage

    movl     8(%esp),%eax            # move pointer to t2 into eax
```

```

    movl    _EAX(%eax),%ebx        # get new value for eax into eb
    movl    %ebx,_eax_save        # save it
    movl    _EBX(%eax),%ebx        # restore old registers
    movl    _ECX(%eax),%ecx
    movl    _EDX(%eax),%edx
    movl    _ESI(%eax),%esi
    movl    _EDI(%eax),%edi
    movl    _EBP(%eax),%ebp
    movl    _ESP(%eax),%esp        # restore stack pointer
    movl    _PC(%eax),%eax         # restore return address into e
    movl    %eax,4(%esp)           # copy over the ret address on
    movl    _eax_save,%eax

    ret

#endif // x86

```

That's where all the definition located.

```

#ifdef x86

/* the offsets of the registers from the beginning of the thread object */
#define _ESP    0
#define _EAX    4
#define _EBX    8
#define _ECX    12
#define _EDX    16
#define _EBP    20
#define _ESI    24
#define _EDI    28
#define _PC     32

/* These definitions are used in Thread::AllocateStack(). */
#define PCState      (_PC/4-1)
#define FPState      (_EBP/4-1)
#define InitialPCState (_ESI/4-1)
#define InitialArgState (_EDX/4-1)
#define WhenDonePCState (_EDI/4-1)
#define StartupPCState (_ECX/4-1)

#define InitialPC      %esi
#define InitialArg      %edx
#define WhenDonePC     %edi
#define StartupPC       %ecx

#endif // x86

```


In ThreadRoot:

Will call the initial functions defined in switch.h(by setting registers).

In SWITCH:

```
movl    %eax, _eax_save      # save the value of eax
movl    4(%esp), %eax        # move pointer to t1 into eax
movl    %ebx, _EBX(%eax)     # save registers
movl    %ecx, _ECX(%eax)
movl    %edx, _EDX(%eax)
movl    %esi, _ESI(%eax)
movl    %edi, _EDI(%eax)
movl    %ebp, _EBP(%eax)
movl    %esp, _ESP(%eax)     # save stack pointer
movl    _eax_save, %ebx      # get the saved value of eax
movl    %ebx, EAX(%eax)      # store it
```

Save the current eax on the stack so that we can use it as a pointer to t1. Then saves the values of the other registers(from t1) to the stack, as well as the stack pointer and the saved value of eax.

```
movl    0(%esp), %ebx        # get return address from stack into ebx
movl    %ebx, _PC(%eax)      # save it into the pc storage

movl    8(%esp), %eax        # move pointer to t2 into eax

movl    _EAX(%eax), %ebx     # get new value for eax into ebx
movl    %ebx, _eax_save      # save it
movl    _EBX(%eax), %ebx     # restore old registers
```

Get the return address from stack and save it in to the pc storage(t1), then move the pointer to t2 into eax, restore the saved value of eax into eax.

```
movl    _EBX(%eax), %ebx     # restore old registers
movl    _ECX(%eax), %ecx
movl    _EDX(%eax), %edx
movl    _ESI(%eax), %esi
movl    _EDI(%eax), %edi
movl    _EBP(%eax), %ebp
movl    _ESP(%eax), %esp     # restore stack pointer
```

Restore the value of old registers and the stack pointer(t2).

```

movl    _PC(%eax),%eax        # restore return address into eax
movl    %eax,4(%esp)          # copy over the ret address on the stack
movl    _eax_save,%eax

ret

```

Restore the return address and copy it over the return address on stack. Finally, return from function, finish the thread switch, start executing t2.

This complete the context switch.

(depends on the previous process state, e.g., [New,Running,Waiting] → Ready)

New -> Ready: 1-1

Running -> Ready: 1-2

Waiting -> Ready: 1-4

For 1-2(Running -> Ready) and 1-4(Waiting -> Ready), they turn to Ready because of Context Switch, so while they Ready -> Running, we have to restore their value, register and PC... and continue from previous place. For 1-1(New -> Ready), when Ready -> Running, just simply start from scratch.

for loop in Machine::Run()

Explained in MP1.

Executing the program.

Implementation

code/lib/debug.h

By the requirement 2-3 in spec, we add a debugging flag **z**.

```

22  const char dbgAll = '+';        // turn on all debug messages
23  const char dbgThread = 't';     // threads
24  const char dbgSynch = 's';      // locks, semaphores, condition vars
25  const char dbgInt = 'i';        // interrupt emulation
26  const char dbgMach = 'm';       // machine emulation
27  const char dbgDisk = 'd';       // disk emulation
28  const char dbgFile = 'f';       // file system
29  const char dbgAddr = 'a';       // address spaces
30  const char dbgNet = 'n';        // network emulation
31  const char dbgSys = 'u';         // systemcall
32  const char dbgTraCode = 'c';
33  const char dbgMP3 = 'z';        // MP3 add

```

code/threads/thread.h

1. Adding `aging(int)` function to be implemented in `thread.cc`.
2. Adding five variables and its `get` and `set` functions to record:
 - burst time
 - accumulated execution time
 - priority of the thread
 - time when the thread enters the ready state
 - time when the thread enters the running state

```
134 // MP3 Add
135 private:
136     double burstTime;
137     double accumulatedExecutionTime;
138     int execPriority;
139     int enterReadyTime;
140     int enterRunningTime;
141 public:
142     void aging(int currentTime);
143     void setBurstTime(double t) { burstTime = t; }
144     double getBurstTime() { return burstTime; }
145     void setAccumulatedExecutionTime(int t) { accumulatedExecutionTime = t; }
146     double getAccumulatedExecutionTime() { return accumulatedExecutionTime; }
147     void setPriority(int priority) { execPriority = priority; }
148     int getPriority() { return execPriority; }
149     void setReadyTick(int t) { enterReadyTime = t; }
150     int getReadyTick() { return enterReadyTime; }
151     void setRunTick(int t) { enterRunningTime = t; }
152     int getRunTick() { return enterRunningTime; }
```

code/threads/thread.cc

1. `Thread::Thread()`
Initialize the five variables we have created.

```
49     burstTime = accumulatedExecutionTime = 0; // MP3 add
50     execPriority = enterReadyTime = enterRunningTime = 0; // MP3 add
```

2. `Thread::Yield()`
Calculate the elapsed time between the current time and when the thread enters the running state, then add the elapsed time to the accumulated

execution time.

```
213     if (nextThread != NULL) {
214         int elapsedTicks = kernel -> stats -> totalTicks - kernel -> currentThread -> getRunTick(); // MP3 add
215         int totalAccumulatedTime = kernel -> currentThread -> getAccumulatedExecutionTime() + elapsedTicks; // MP3 add
216         kernel -> currentThread -> setAccumulatedExecutionTime(totalAccumulatedTime); // MP3 add
217         DEBUG(dbgMP3, "[E] Tick " << kernel->stats->totalTicks << // MP3 add
218             "Thread " << nextThread->getID() << "is now selected for execution, thread " << // MP3 add
219             kernel->currentThread->ID << " is replaced, and it has executed" << // MP3 add
220             kernel->currentThread->getAccumulatedExecutionTime() << "ticks \n" ); // MP3 add
221         kernel->scheduler->ReadyToRun(this);
222         kernel->scheduler->Run(nextThread, FALSE);
223     }
```

3. Thread::Sleep()

Update the accumulated execution time like that in Thread::Yield(), then calculate the approximated burst time given by spec and set the burst time to the current thread.

```
259     int elapsedTicks = kernel -> stats -> totalTicks - kernel -> currentThread -> getRunTick(); // MP3 add
260     int totalAccumulatedTime = kernel -> currentThread -> getAccumulatedExecutionTime() + elapsedTicks; // MP3 add
261     kernel -> currentThread -> setAccumulatedExecutionTime(totalAccumulatedTime); // MP3 add
262     DEBUG(dbgMP3, "[D] Tick [" << kernel -> stats -> totalTicks << // MP3 add
263         "]: Thread [" << kernel -> currentThread -> getID() << // MP3 add
264         "]" update approximate burst time, from: ["<< kernel -> currentThread -> getBurstTime() << // MP3 add
265         "], add [" << kernel -> currentThread -> getAccumulatedExecutionTime() << "], to [" << // MP3 add
266         double(kernel -> currentThread -> getBurstTime() * 0.5) + // MP3 add
267         double(kernel -> currentThread -> getAccumulatedExecutionTime() * 0.5) << "]" \n"); // MP3 add
268     double newBurstTime = 0.5 * kernel -> currentThread -> getBurstTime() + // MP3 add
269         0.5 * kernel -> currentThread -> getAccumulatedExecutionTime(); // MP3 add
270     kernel -> currentThread -> setBurstTime(newBurstTime); // MP3 add
271     kernel -> currentThread -> setAccumulatedExecutionTime(0); // MP3 add
```

4. Thread::aging()

Calculate the waiting time and update the priority while waiting for more than 1500 ticks.

```
460     void Thread::aging(int currentTime) { // MP3 add
461         int prevPriority = execPriority;
462         int newPriority = execPriority;
463         int waitingTime = currentTime - enterReadyTime;
464
465         while(waitingTime > 1500) {
466             newPriority += 10;
467             waitingTime -= 1500;
468             enterReadyTime += 1500;
469         }
470         newPriority = min(newPriority, 149);
471         if (prevPriority != newPriority)
472             DEBUG(dbgMP3, "[C] Tick [" << currentTime << "]: Thread [" << ID <<
473                 "]" changes its priority from [" << prevPriority << "]" to [" << newPriority << "]\n");
474         execPriority = newPriority;
475     }
```

code/threads/kernel.h

1. Modify the Exec function so that we can also send the thread priority.

```
// int Exec(char* name); // MP3 Delete
int Exec(char *name, int priorityNumber); // MP3 Add
```

2. Add `threadPriorityNumber` since each thread has its priority.

```
75     private:
76
77     Thread* t[10];
78     char*   execfile[10];
79     int threadPriorityNumber[10]; // MP3 Add
80     int execfileNum;
```

code/threads/kernel.cc

1. Add the flag `-ep` like the flag `-e`, and initialize each `threadPriorityNumber`.

```
79     } else if (strcmp(argv[i], "-ep") == 0) { // MP3 add
80         execfile[++execfileNum] = argv[++i];
81         threadPriorityNumber[execfileNum] = atoi(argv[++i]);
82         cout << execfile[execfileNum] << " has the priority of " << threadPriorityNumber[execfileNum] << "\n";
83         // check whether the priority is valid
84         ASSERT(0 <= threadPriorityNumber[execfileNum] && threadPriorityNumber[execfileNum] <= 149);
85     }
```

2. Add `t[threadNum] -> setPriority(priorityNumber)` to each thread, and delete the original `Exec` function since it won't be used anymore.

```
272     int Kernel::Exec(char *name, int priorityNumber) { // MP3 Add
273         t[threadNum] = new Thread(name, threadNum);
274         t[threadNum] -> setPriority(priorityNumber);
275         t[threadNum] -> space = new AddrSpace();
276         t[threadNum] -> Fork((VoidFunctionPtr) &ForkExecute, (void *)t[threadNum]);
277
278         ++threadNum;
279
280         return threadNum - 1;
281     }
282
283     // int Kernel::Exec(char* name) // MP3 Delete
284     // {
285     //     t[threadNum] = new Thread(name, threadNum);
286     //     t[threadNum]->space = new AddrSpace();
287     //     t[threadNum]->Fork((VoidFunctionPtr) &ForkExecute, (void *)t[threadNum]);
288     //     threadNum++;
289
290     //     return threadNum-1;
```

3. Modify the calling of `Exec()` in `ExecAll()` .

```
269 void Kernel::ExecAll()
270 {
271     for (int i=1;i<=execfileNum;i++) {
272         int a = Exec(execfile[i], threadPriorityNumber[i]); // MP3 Modify
273     }
274     currentThread->Finish();
275     //Kernel::Exec();
276 }
```

code/threads/scheduler.h

1. Use `SortedList` to create `L1` and `L2` , and just use a simple list for `L3` since `L1` and `L2` have their own sorting rule. Remove the original `readyList` at the same time.

```
48 private:
49 /* MP3 delete
50 List<Thread *> *readyList; // queue of threads that are ready to run, but not running
51 */
52 SortedList<Thread *> *L1, *L2; // MP3 add
53 List<Thread *> *L3; // MP3 add
54 Thread *toBeDestroyed; // finishing thread to be destroyed
55 // by the next thread that runs
```

2. Add three functions: `UpdatePriority()` , `Scheduling()` , `SortUpdatedList()` to be implemented in `scheduler.cc` , and add query functions of the three lists for `alarm.cc` to call. The detail will be explained in `alarm.cc` .

```
36 void UpdatePriority(); // MP3 add
37 void Scheduling(); // MP3 add
38 void SortUpdatedList(); // MP3 add
39
40 bool isL1Empty() { return L1 -> IsEmpty(); } // MP3 add
41 bool isL2Empty() { return L2 -> IsEmpty(); } // MP3 add
42 bool isL3Empty() { return L3 -> IsEmpty(); } // MP3 add
43 Thread* L1Front() { return L1 -> Front(); } // MP3 add
44 Thread* L2Front() { return L2 -> Front(); } // MP3 add
45 Thread* L3Front() { return L3 -> Front(); } // MP3 add
```

code/threads/scheduler.cc

1. Comparison function for L1 . Calculate the remaining burst time and return the priority.

```
26 // MP3 add
27 int PreemptiveSJF(Thread* First, Thread* Second) {
28     double remain_burst_time1 = (First -> getBurstTime()) - (First -> getAccumulatedExecutionTime());
29     double remain_burst_time2 = (Second -> getBurstTime()) - (Second -> getAccumulatedExecutionTime());
30     if (remain_burst_time1 < remain_burst_time2)
31         return -1;
32     else if (remain_burst_time1 == remain_burst_time2)
33         return 0;
34     else
35         return 1;
36 }
```

2. Comparison function for L2 . Just compare its priority.

```
38 // MP3 add
39 int NonPreemptive(Thread *First, Thread *Second) {
40     if (First -> getPriority() > Second -> getPriority())
41         return -1;
42     else if (First -> getPriority() == Second -> getPriority())
43         return 0;
44     else
45         return 1;
46 }
```

3. Modify the constructor

```
54 Scheduler::Scheduler()
55 {
56     // readyList = new List<Thread *>; // MP3 delete
57     L1 = new SortedList<Thread *>(PreemptiveSJF); // MP3 add
58     L2 = new SortedList<Thread *>(NonPreemptive); // MP3 add
59     L3 = new List<Thread *>; // MP3 add
60     toBeDestroyed = NULL;
61 }
```

4. Modify the destructor

```
68 Scheduler::~~Scheduler()
69 {
70     // delete readyList; // MP3 delete
71     delete L1, L2, L3; // MP3 add
72 }
```


5. UpdatePriority : If the list is not empty, we then iterate all its threads and call aging.

```
74 // MP3 add
75 void Scheduler::UpdatePriority() {
76     int ticks = kernel -> stats -> totalTicks;
77
78     if (!isL1Empty()) {
79         for (ListIterator<Thread*> *iter = new ListIterator<Thread*>(L1); !iter -> IsDone(); iter -> Next())
80             iter -> Item() -> aging(ticks);
81     }
82
83     if (!isL2Empty()) {
84         for (ListIterator<Thread*> *iter = new ListIterator<Thread*>(L2); !iter -> IsDone(); iter -> Next())
85             iter -> Item() -> aging(ticks);
86     }
87
88     if (!isL3Empty()) {
89         for (ListIterator<Thread*> *iter = new ListIterator<Thread*>(L3); !iter -> IsDone(); iter -> Next())
90             iter -> Item() -> aging(ticks);
91     }
92
93     return;
94 }
```

6. Scheduling : For each thread in L2 , if its priority is greater than 99, then move it to L1 . For each thread in L3 , if its priority is greater than 99 or 49, then move it to L1 or L2 , respectively. Once a thread is removed from or insert to a list, print a debug message.

```
96 // MP3 add
97 void Scheduler::Scheduling() {
98     ListIterator<Thread*> *iter;
99     Thread* currentThread;
100     if (!isL2Empty()) {
101         for (ListIterator<Thread*> *iter = new ListIterator<Thread*>(L2); !iter -> IsDone(); iter -> Next()) {
102             currentThread = iter -> Item();
103             int threadPriority = iter -> Item() -> getPriority();
104             if (threadPriority > 99) { // put currentThread from L2 to L1
105                 L2 -> Remove(currentThread);
106                 DEBUG(dbgMP3, "[B] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
107                     "]" is removed from queue L[2]\n");
108                 L1 -> Insert(currentThread);
109                 DEBUG(dbgMP3, "[A] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
110                     "]" is inserted into queue L[1]\n");
111             }
112         }
113     }
114
115     if (!isL3Empty()) {
116         for (ListIterator<Thread*> *iter = new ListIterator<Thread*>(L3); !iter -> IsDone(); iter -> Next()) {
117             currentThread = iter -> Item();
118             int threadPriority = iter -> Item() -> getPriority();
119             if (threadPriority > 99) { // put currentThread from L3 to L1
120                 L3 -> Remove(currentThread);
121                 DEBUG(dbgMP3, "[B] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
122                     "]" is removed from queue L[3]\n");
123                 L1 -> Insert(currentThread);
124                 DEBUG(dbgMP3, "[A] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
125                     "]" is inserted into queue L[1]\n");
126             } else if (threadPriority > 49) { // put currentThread from L3 to L2
127                 L3 -> Remove(currentThread);
128                 DEBUG(dbgMP3, "[B] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
129                     "]" is removed from queue L[3]\n");
130                 L2 -> Insert(currentThread);
131                 DEBUG(dbgMP3, "[A] Tick[" << kernel -> stats -> totalTicks << "]: Thread [" << currentThread -> getID() <<
132                     "]" is inserted into queue L[2]\n");
133             }
134         }
135     }
136 }
```


7. Before scheduling the next thread onto the CPU, we call `SortUpdatedList` to ensure the priority in `L1` and `L2` are well sorted.

```
138 // MP3 add
139 void Scheduler::SortUpdatedList() {
140     SortedList<Thread *> *tmpL1 = new SortedList<Thread *>(PreemptiveSJF);
141     SortedList<Thread *> *tmpL2 = new SortedList<Thread *>(NonPreemptive);
142     if (!isL1Empty()) {
143         for (ListIterator<Thread *> *iter = new ListIterator<Thread *>(L1); !iter -> IsDone(); iter -> Next())
144             tmpL1 -> Insert(iter -> Item());
145         delete L1;
146         L1 = tmpL1;
147     }
148
149     if (!isL2Empty()) {
150         for (ListIterator<Thread *> *iter = new ListIterator<Thread *>(L2); !iter -> IsDone(); iter -> Next())
151             tmpL2 -> Insert(iter -> Item());
152         delete L2;
153         L2 = tmpL2;
154     }
155     return;
156 }
```

8. Modify `ReadyToRun`. Put the thread into the corresponding Ready queue.

```
166 void
167 Scheduler::ReadyToRun (Thread *thread)
168 {
169     ASSERT(kernel->interrupt->getLevel() == IntOff);
170     DEBUG(dbgThread, "Putting thread on ready list: " << thread->getName());
171     //cout << "Putting thread on ready list: " << thread->getName() << endl ;
172     // thread->setStatus(READY); // MP3 delete
173     // readyList->Append(thread); // MP3 delete
174     // MP3 add
175     thread -> setReadyTick(kernel -> stats -> totalTicks);
176     int threadPriority = thread -> getPriority();
177     if (threadPriority < 50) {
178         L3 -> Append(thread);
179         DEBUG(dbgMP3, "[A] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << thread -> getID() <<
180             "]" is inserted into queue L[3]\n");
181     } else if (threadPriority < 100) {
182         L2 -> Insert(thread);
183         DEBUG(dbgMP3, "[A] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << thread -> getID() <<
184             "]" is inserted into queue L[2]\n");
185     } else {
186         L1 -> Insert(thread);
187         DEBUG(dbgMP3, "[A] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << thread -> getID() <<
188             "]" is inserted into queue L[1]\n");
189     }
190     return;
191 }
```

9. Calling `SortUpdateList()` in the begining. Check `L1`, `L2`, and `L3` to see if there are threads. Since we have sort the lists, the front thread of the lists has the highest priority. Once a thread is found, schedule it onto the CPU and remove

it from the list. If no threads are found, return NULL .

```
201 Thread *
202 Scheduler::FindNextToRun ()
203 {
204     ASSERT(kernel->interrupt->getLevel() == IntOff);
205
206     // MP3 delete
207     /*if (readyList->IsEmpty()) {
208         return NULL;
209     } else {
210         return readyList->RemoveFront();
211     }*/
212     // MP3 add
213     SortUpdatedList();
214     Thread* nextThread = NULL;
215     if (!isL1Empty()) {
216         nextThread = L1 -> RemoveFront();
217         nextThread -> setRunTick(kernel -> stats -> totalTicks);
218         DEBUG(dbgMP3, "[B] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << nextThread -> getID() <<
219             "]" is removed from queue L[1]\n");
220     } else if (!isL2Empty()) {
221         nextThread = L2 -> RemoveFront();
222         nextThread -> setRunTick(kernel -> stats -> totalTicks);
223         DEBUG(dbgMP3, "[B] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << nextThread -> getID() <<
224             "]" is removed from queue L[2]\n");
225     } else if (!isL3Empty()) {
226         nextThread = L3 -> RemoveFront();
227         nextThread -> setRunTick(kernel -> stats -> totalTicks);
228         DEBUG(dbgMP3, "[B] Tick [" << kernel -> stats -> totalTicks << "]: Thread [" << nextThread -> getID() <<
229             "]" is removed from queue L[3]\n");
230     }
231     return nextThread;
232 }
```

10. In Run() , once the old thread is switched back, we print the debug message and update its enterRunningTime .

```
281 SWITCH(oldThread, nextThread);
282
283 DEBUG(dbgMP3 , "[E] Tick [" << kernel -> stats -> totalTicks <<
284     "]" : Thread ["<< nextThread -> getID() << "]" is now selected for execution, thread [" <<
285     oldThread -> getID() << "]" is replaced, and it has executed [" <<
286     kernel -> stats -> totalTicks - oldThread -> getRunTick() << "]" ticks \n");
287
288 // we're back, running oldThread
289 oldThread -> setRunTick(kernel -> stats -> totalTicks);
```

code/threads/alarm.cc

In the begining, we update the lists by calling UpdatePriority() , Scheduling() , and SortUpdatedList() . Then we add the following rules to the Callback() function to decide whether the current thread should be preempted.

- The priority of the current thread is greater than 99, and the remain burst time of the front thread of L1 is smaller than that of the current thread.
- The priority of the current thread is greater than 49 but less than 100, and L1 has thread.
- The priority of the current thread is less than 50 and there exists a thread in L1 , L2 , or L3 .

Since we can not access the private data L1 , L2 , and L3 of the scheduler class, we implement the public functions in scheduler.h .

```
46 void
47 Alarm::CallBack()
48 {
49     Interrupt *interrupt = kernel->interrupt;
50     MachineStatus status = interrupt->getStatus();
51
52     if (status != IdleMode) {
53         // interrupt->YieldOnReturn(); // MP3 delete
54         kernel -> scheduler -> UpdatePriority();
55         kernel -> scheduler -> Scheduling();
56         kernel -> scheduler -> SortUpdatedList();
57         if (kernel -> currentThread -> getPriority() > 99) {
58             if (!kernel -> scheduler -> isL1Empty()) {
59                 Thread* L1Front = kernel -> scheduler -> L1Front();
60                 Thread* currentThread = kernel -> currentThread;
61                 double remainBurstTime1 = currentThread -> getBurstTime() - currentThread -> getAccumulatedExecutionTime();
62                 double remainBurstTime2 = L1Front -> getBurstTime() - L1Front -> getAccumulatedExecutionTime();
63                 if (remainBurstTime1 > remainBurstTime2)
64                     interrupt -> YieldOnReturn();
65             }
66         } else if (kernel -> currentThread -> getPriority() > 49) {
67             if (!kernel -> scheduler -> isL1Empty())
68                 interrupt -> YieldOnReturn();
69         } else {
70             if ((!kernel -> scheduler -> isL1Empty()
71                 || !kernel -> scheduler -> isL2Empty()
72                 || !kernel -> scheduler -> isL3Empty()))
73                 interrupt -> YieldOnReturn();
74         }
75     }
76 }
```